

COMS 30115

Carl Henrik Ek - carlhenrik.ek@bristol.ac.uk

February 18th, 2019



<http://www.carlhenrik.com>

- Shading/Lighting equations
- BRDFs
 - a more realistic way to parametrise ray-surface interaction
- Fun things to do with normals
 - Bump and Environment mapping

Today

- Shadows
 - Last element of raytracer lab
- Cameras
- Optimisations
 - ideas
- Summarise raytracing

Where are we



- Shadows [URL](#)
- Cameras [URL](#)
- Optimisations
 - [Jacco Bikker Flipcode](#)
 - [Gavan Woolery, Gamasutra, Why I still Think Raytracing is the Future](#)
 - [Accelerated Structures, Scratchapixel](#)
 - [Robin Marcus, Realtime rendering blog](#) - 4 part article about realtime raytracing

Shadows


Shadows




1

¹X-Files Universe


Raytracing




Raytracing




Raytracing




Raytracing



Raytracing



Raytracing




Shadow Acne

- Self-intersection
- If our primitive is a triangle we do not have self-shadows
- *you will see this in the lab*
- shadow bias




Shadow Acne


- Self-intersection
- If our primitive is a triangle we do not have self-shadows
- *you will see this in the lab*
- shadow bias



Shadow Acne

- Self-intersection
- If our primitive is a triangle we do not have self-shadows
- *you will see this in the lab*
- shadow bias






Code

```
/*compute primary ray*/
for(int i=0;i<N_primitives;i++)
    {/*compute intersection*/}
/*compute shadow ray*/
for(int i=0;i<N_primitives;i++)
    {/*compute intersection*/}
glm::vec3 i_tot = glm::vec3(0,0,0);
for(int i=0;i<N_lights;i++)
    {/*compute light*/
        i_tot += ;
    }
```


Soft Shadows



2

²[Image URL](#)


Soft Shadows



²[Image URL](#)


Soft Shadows

- Zero dimensional things aren't really that common
- Most light-sources have spatial extent
- **umbra** - hard shadow
- **penumbra** - partial shadow
- **antumbra** - cross over




Soft Shadows

- Zero dimensional things aren't really that common
- Most light-sources have spatial extent
- **umbra** - hard shadow
- **penumbra** - partial shadow
- **antumbra** - cross over
- Approximate using several point/spot lights



Soft Shadows

- Zero dimensional things aren't really that common
- Most light-sources have spatial extent
- **umbra** - hard shadow
- **penumbra** - partial shadow
- **antumbra** - cross over
- Approximate using several point/spot lights
- Approximate with several intersection planes




- Shadows adds a lot of realism
- Expensive to compute as you will need more rays
- If you use hard shadows compensate with a lot of ambient light

Cameras




Pinhole Cameras




- Raytraced images often looks too "clean"
- infinite depth-of-field

Anti-Aliasing




- shoot several rays per pixel and blend colours together

- Film needs to be bombarded with sufficient energy to generate a colour
- Opening for rays bigger than a one photon wide pinhole: **Aperture**
- Same as glossy reflection, each ray hitting the eye comes from different parts of the world




Depth-of-Field³




³<https://www.bhphotovideo.com/explora/photography/tips-and-solutions/how-focus-works>

Depth-of-Field³




³<https://www.bhphotovideo.com/explora/photography/tips-and-solutions/how-focus-works>

Depth-of-Field³



³<https://www.bhphotovideo.com/explora/photography/tips-and-solutions/how-focus-works>

Depth-of-Field



- Which rays to trace, how many?
- Different geometry aperture creates different images



SIX BLADED APERTURE



SEVEN BLADED APERTURE

Lens Flare⁴

- A real camera has a lens
- Transition between two media causes refraction
- Light bounces inside lens
- Refraction calculations inside lens




⁴[URL](#)

- We have seen quite a few things now
 - reflection, refraction
 - shadows
 - depth-of-field
 - reflection models
 - material properties
 - etc.

Summary


- We have seen quite a few things now
 - reflection, refraction
 - shadows
 - depth-of-field
 - reflection models
 - material properties
 - etc.
- All follows the same principle
 - *basic physics*
 - *If light can be assumed to be a ray, then we just have to follow it along its path!*

Random is your friend⁵



- *We (humans) are really good at picking up regularities, avoid them for increased fidelity*

⁵<http://www-alg.ist.hokudai.ac.jp/~jan/randsphere.pdf>



⁶https://fabiansanglard.net/rayTracing_back_of_business_card/index.php



⁷https://fabiansanglard.net/postcard_pathtracer/

Optimisation

Ek In the first part of the course we will look at raytracing

Ek In the first part of the course we will look at raytracing

UoB Damn this is slow, I can't even move the camera, or wait, did it move?

Ek In the first part of the course we will look at raytracing

UoB Damn this is slow, I can't even move the camera, or wait, did it move?

Ek Well thats ray-tracing, try to reduce the screen-size to 10×10 so that we can see something


Ek In the first part of the course we will look at raytracing

UoB Damn this is slow, I can't even move the camera, or wait, did it move?

Ek Well thats ray-tracing, try to reduce the screen-size to 10×10 so that we can see something

UoB Really useful this unit is, the graphics is both ugly and slow


Smash of Fairlight







5 Faces

- Refraction
- Secondary rays in several levels (lots)
- Depth of Field
- Implicit surfaces
- Real-Time




- Structure our data to do less calculations
 - Bounding Boxes, Tree's, etc.
- Optimise your code
- Reduce visual quality/correctness for speed



Bounding Boxes

- Intersections are expensive to compute
- Create simpler geometry that surrounds object
- The tighter the bounding volume the better but that usually requires more vertices



- Do we need to check all object everywhere in space?
- Create a KD-Tree that splits up space
- Find large areas with single object

Cramer's Rule



$$\begin{aligned}\begin{pmatrix} t \\ u \\ v \end{pmatrix} &= \frac{1}{\det(-\mathbf{d}, -\mathbf{e}_1, -\mathbf{e}_2)} \begin{pmatrix} \det(-\mathbf{s}, -\mathbf{e}_1, -\mathbf{e}_2) \\ \det(-\mathbf{d}, -\mathbf{s}, -\mathbf{e}_2) \\ \det(-\mathbf{d}, -\mathbf{e}_1, -\mathbf{s}) \end{pmatrix} \\ &= \{ \det(\mathbf{a}, \mathbf{b}, \mathbf{c}) = -(\mathbf{a} \times \mathbf{c})^T \mathbf{b} = -(\mathbf{c} \times \mathbf{b})^T \mathbf{a} \} \\ &= \frac{1}{\mathbf{p}^T \mathbf{e}_1} = \begin{pmatrix} \mathbf{q}^T \mathbf{e}_2 \\ \mathbf{p}^T \mathbf{s} \\ \mathbf{q}^T \mathbf{d} \end{pmatrix}\end{aligned}$$

$$\mathbf{p} = \mathbf{d} \times \mathbf{e}_2$$

$$\mathbf{q} = \mathbf{s} \times \mathbf{e}_1$$

- Computing an inverse is more work $\mathbf{A}^T \mathbf{A} = \mathbf{I}$

Refraction Rays



Refraction

$$\mathbf{t} = \left(\frac{n_i}{n_t} \cos(\theta_i) - \sqrt{1 - \left(\frac{n_i}{n_t} \right)^2 \cos^2(\theta_i)} \right) \mathbf{n} - \frac{n_i}{n_t} \mathbf{i}$$

Refraction Rays

Code

```
float b = Ni/Nt; float b2 = b*b;
float ni =N.x*I.x+N.y*I.y+N.z*I.z; //3m+2a
float ni2 = ni*ni; //1m
float D2 = 1.0f - b2*ni2; //1m+1a
a = b*ni-sqrtf(D2); //1m+1a+1sqrt
T.x = a*N.x-b*I.x;
T.y = a*N.y-b*I.y;
T.z = a*N.z-b*I.z; //6m+3a
```

Total: 12 multiplications, 6 additions, 1 square-root

⁸<http://hugi.scene.org/online/hugi23/torefrac.htm> Code on the course website

Code Optimisation

- What can we pre-compute?
- Something that does not have a massive range
- Something that has few indexing variables

⁸<http://hugi.scene.org/online/hugi23/torefrac.htm> Code on the course website

Refraction Rays

a is just a function of **ni** it can be pre-computed and tabled (**ni** is bounded)

Code

```
float scalefac = 16384;
float *aLUT;

float ni=N.x*I.x+N.y*I.y+N.z*I.z; //3m+2a
float a=aLUT[(int)(ni*scalefac)]; //1m+1lu

T.x = a*N.x-b*I.x;
T.y = a*N.y-b*I.y;
T.z = a*N.z-b*I.z; //6m+3a
```

Total: 10 multiplications, 5 additions, 1 look-up

⁸<http://hugi.scene.org/online/hugi23/torefrac.htm> Code on the

Refraction Rays

Do we need the vector to be normalised, if not we can take $g = \frac{a}{b}$

Code

```
float scalefac = 16384;
float *gLUT; //(a/b)

float ni=N.x*I.x+N.y*I.y+N.z*I.z; //3m+2a
float g=gLUT[(int)(ni*scalefac)]; //1m+1lu

T.x = g*N.x-I.x;
T.y = g*N.y-I.y;
T.z = g*N.z-I.z; //3m+3a
```

- Total: 7 multiplications, 5 additions, 1 look-up

- *This version is about twice as fast as the first⁸*

Summary

Summary

- Calculate number of operations
 - is there a less general way to do this
 - are there special cases, then code all of them
- Use heuristics
 - what is the most general case I'll see
- What shows and what doesn't?
- Profile your code

- Know your two weapons

- Inner product

$$\mathbf{x}^T \mathbf{y} = ||\mathbf{x}|| ||\mathbf{y}|| \cos(\theta)$$

- Outer product

$$\mathbf{x} \times \mathbf{y} = \mathbf{z}$$

$$\mathbf{z} \perp \mathbf{x}$$

$$\mathbf{z} \perp \mathbf{y}$$

- Think how things work physically

- how can we "mimick" this behaviour?

- *Thats rendering for you*

Next Time

Lecture Friday 22nd of February

- Simon will talk about generating geometry

Lab Continue with Lab 1 and try to finish up to 50%

eof